Open in app ↗

Palantir's Real Secret Sauce — Ontologies

4 min read · Mar 25, 2026

balaji bal Follow

Listen Share ••• More

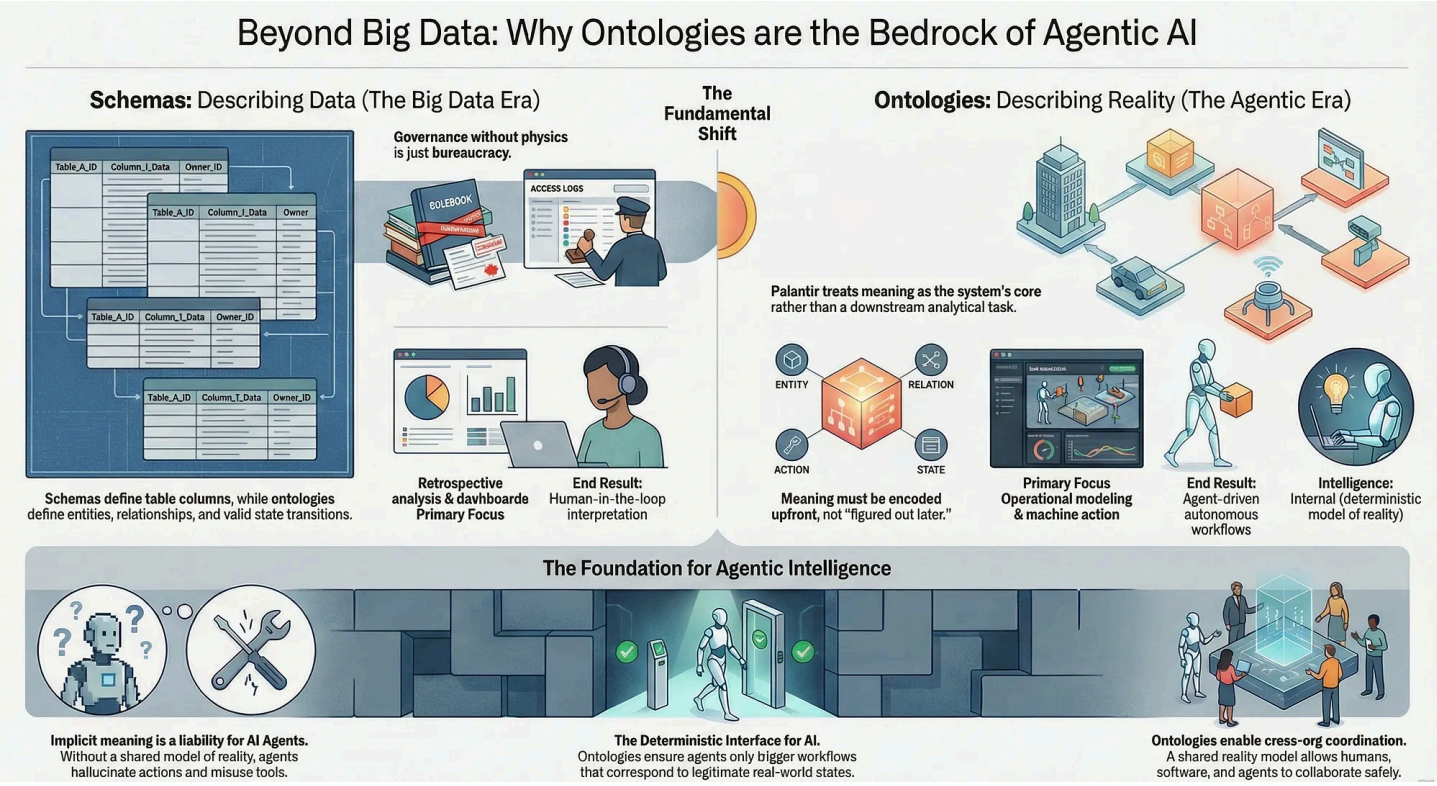

Palantir distinguished itself by building ontologies that model real-world entities and their relationships.

For more than a decade, the big data industry has obsessed over the same set of problems: data quality, data governance, lineage, catalogs, ownership. Entire product categories and consulting empires have been built around answering questions like *Who owns this data?*, *Where did it come from?*, and *Can we trust it?*

These questions matter. But they are not the whole story.

What the industry mostly missed until recently is that clean, governed data is still inert unless a system understands **what the data actually represents**. Not just structurally, but *semantically*. Not just rows and columns, but *things in the world*, how they relate, and what can happen to them.

This is where ontologies come in. And this is why Palantir, often misunderstood and sometimes dismissed, quietly built one of the most important software moats of the last twenty years.

Palantir's Unfashionable Bet

If you read the article *"Understanding Palantir: Forward-Deployed Engineers and the Making of an Unusual Platform Company"*, one thing becomes obvious: Palantir did not start as a dashboards company, a data lake company, or even an analytics company. From the beginning, it was a **world-modeling company**.

While the rest of the big data ecosystem was racing to ingest more data faster, Palantir spent an unusual amount of time asking a different question:

> What are the actual entities in this domain, and how do they relate?

Their answer was the ontology.

This was an unfashionable bet. Ontologies sounded academic, slow, and expensive at a time when Hadoop clusters and schema-on-read promised speed and flexibility. Most vendors treated meaning as something analysts or downstream applications would "figure out later."

> Palantir did the opposite. They insisted that meaning had to be encoded up front, explicitly, and shared across the system.

That decision shaped everything that followed.

What Palantir Means by "Ontology" (And What It Doesn't)

In Palantir's world, an ontology is not a philosophical taxonomy or a metadata glossary. It is an **operational model of reality**.

Concretely, it defines:

- **Entities** (people, accounts, vehicles, facilities, transactions)
- **Relationships** (ownership, movement, dependency, authorization)
- **Constraints** (what is allowed, what is forbidden)
- **State transitions** (how things change over time)

This is the distinction the industry largely missed:

> Schemas describe data. Ontologies describe reality.

A schema tells you that a table has a `status` column.
An ontology tells you what *states are possible*, how you move between them, and what actions those states enable or prohibit.

Once you see this, Palantir's architecture makes sense. The ontology isn't a layer on top of the system — it *is* the system.

Why Big Data Used to Work Without Ontologies

It's important to be fair to the past.

Early big data systems *did* work — at least well enough — without ontologies. Hadoop, Spark, and data lakes succeeded because they optimized for a specific mode of computing:

- Retrospective analysis
- Batch processing
- Human-in-the-loop interpretation

In these systems, meaning lived mostly in people's heads. Analysts joined tables, interpreted results, and made decisions outside the system. Data governance filled the gaps with process, policy, and documentation.

> You could get away with this because the system itself wasn't expected to act. It produced insight;

> humans supplied judgment.

Ontologies were optional because intelligence was external.

The Breaking Point: Agentic Systems

That era is ending.

Agentic systems don't just analyze data, they **take actions**. They call tools, trigger workflows, allocate resources, and make decisions at machine speed. The moment you move from insight to action, implicit meaning becomes a liability.

Without ontologies:

- Agents hallucinate actions that don't make sense
- Tools are misused because their effects aren't well-defined
- Safety boundaries become fuzzy or unenforceable
- Coordination between agents collapses into prompt spaghetti

This is why so many "AI agents" today look impressive in demos and fail in production. They are reasoning over text and schemas, not over a shared model of reality.

Palantir's systems don't have this problem — because their agents are constrained by the ontology. Actions are only valid if they correspond to legitimate state transitions on real entities.

This is not a nice-to-have. It's table stakes.

Ontologies as the Coordination Layer

One of the most underappreciated aspects of ontologies is that they enable **coordination**.

When multiple humans, services, and agents operate in the same system, they need a shared understanding of:

- What exists
- What can change

- Who is allowed to change it
- What the consequences are

Ontologies provide that shared substrate. They make systems composable. They turn APIs into contracts about reality, not just function signatures.

This is why Palantir deployments scale across organizations and years, not just workloads. The ontology becomes the stable interface between people, software, and now agents.

> In modern terms: ontologies are the deterministic interface that agentic systems desperately need.

Why Data Governance Alone Will Never Be Enough

Data governance answers questions like:

- Who owns this dataset?
- Who can access it?
- Is it compliant?

Ontologies answer a different class of questions:

- What *is* this thing?
- What can happen to it?
- What actions are valid right now?

Governance without ontology is bureaucracy without physics. You can regulate access to data all you want, but if the system doesn't understand the domain it's operating in, it cannot reason safely or effectively.

Palantir understood this early. That's why governance in their systems is inseparable from the ontology — it's enforced through meaning, not bolted on with policy.

Palantir Was Early, Not Weird

For years, Palantir's emphasis on ontologies made it seem slow, expensive, or over-engineered. In hindsight, it looks prescient.

As the industry moves toward agentic systems, autonomous workflows, and AI that acts rather than advises, the missing layer has become painfully obvious. Clean data is not enough. Governed data is not enough.

Systems need a shared, explicit model of reality.

Palantir's real secret sauce wasn't defense contracts or sales tactics. It was a simple but radical belief: ***Meaning precedes intelligence.***

Ontology | Agentic Ai | AI | Streamzero | Big Data

Follow

Written by balaji bal

549 followers · 514 following

Founder @ HeadGym.com where we are building The AI Workspace. I write mainly about Applied AI, Software Architecture and DevOps.

No responses yet

Andre Banki

What are your thoughts?